

Dynatrace vs. Datadog

Three differentiators for AI engineers · September 2026

1

Seven layers, not just the agent trace

Dynatrace traces from the GPU and vector store up through the orchestration chain, the agent, and all the way to business impact — in one view. Datadog's story centers on the agent trace and tool decisions. If a bad response traces back to a retrieval failure or a cost spike on the infra side, Datadog puts that in a different place. Dynatrace puts it on the same problem card.

Source: corpus/product-truth.md, verified 2026-09-15 · corpus/competitors/datadog.md, verified 2026-08-28

2

Your production traffic is your eval dataset

dt-evals samples live GenAI spans, runs a judge LLM of your choice across 15 built-in evaluators, and writes scores back to the same trace — no separate eval pipeline. You can gate deploys on quality regression the same way you gate on latency. Datadog has offline experimentation and dataset versioning, but it's a separate workflow from production monitoring.

Honest gap: Datadog is ahead on the pre-production developer loop today (golden datasets, human annotation, offline experiments). dt-evals closes the production side.

Source: corpus/product-truth.md, verified 2026-09-15

3

Causal root cause, not correlation

Davis AI uses deterministic, causal analysis — topology context from Smartscape — to tell you why something broke, not just that it did. When an AI app degrades, it reasons across the full dependency graph to pinpoint the cause. Datadog surfaces outliers and patterns; it does not claim causal, deterministic root cause.

Source: corpus/product-truth.md, verified 2026-09-15 · corpus/competitors/datadog.md, verified 2026-08-28

All claims grounded in corpus/ only. Audience: AI engineers / developers. Grounding verified 2026-09-15.